

Problem Statement

Amazon is an online shopping website that now caters to millions of people everywhere. Over 34,000 consumer reviews for Amazon brand products like Kindle, Fire TV Stick and more are provided. The dataset has attributes like brand, categories, primary categories, reviews.title, reviews.text, and the sentiment. Sentiment is a categorical variable with three levels "Positive", "Negative", and "Neutral". For a given unseen data, the sentiment needs to be predicted. You are required to predict Sentiment or Satisfaction of a purchase based on multiple features and review text.

In [2]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

import nltk
from nltk.corpus import stopwords
from nltk.classify import SklearnClassifier

from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt
%matplotlib inline
from subprocess import check_output

from sklearn.feature_extraction.text import CountVectorizer
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
from sklearn.model_selection import train_test_split
from keras.utils.np_utils import to_categorical
import re
```

Using TensorFlow backend.

In [3]:

```
train = pd.read_csv("../input/ecommerce/train_data.csv")
test = pd.read_csv("../input/ecommerce/test_data.csv")
```

In [4]:

```
train.head()
```

Out[4]:

	name	brand	categories	primaryCategories	reviews.date	re
0	All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi...	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	Electronics	2016-12-26T00:00:00.000Z	Pr or Fr - Pr
1	Amazon - Echo Plus w/ Built-In Hub - Silver	Amazon	Amazon Echo,Smart Home,Networking,Home & Tools...	Electronics,Hardware	2018-01-17T00:00:00.000Z	Ip tv A Ec ar de
2	Amazon Echo Show Alexa-enabled Bluetooth Speak...	Amazon	Amazon Echo,Virtual Assistant Speakers,Electro...	Electronics,Hardware	2017-12-20T00:00:00.000Z	Ju av A op D a
3	Fire HD 10 Tablet, 10.1 HD Display, Wi-Fi, 16 ...	Amazon	eBook Readers,Fire Tablets,Electronics Feature...	Office Supplies,Electronics	2017-08-04T00:00:00.000Z	ve pr Es w w ar
4	Brand New Amazon Kindle Fire 16gb 7" Ips Displ...	Amazon	Computers/Tablets & Networking,Tablets & eBook...	Electronics	2017-01-23T00:00:00.000Z	TI 3i pu l'\

In [5]:

```
test.head()
```

Out[5]:

	name	brand	categories	primaryCategories	reviews.date
0	Fire Tablet, 7 Display, Wi-Fi, 16 GB - Include...	Amazon	Fire Tablets,Computers/Tablets & Networking,Ta...	Electronics	2016-05-23T00:00:00.000Z
1	Amazon Echo Show Alexa-enabled Bluetooth Speake...	Amazon	Computers,Amazon Echo,Virtual Assistant Speake...	Electronics,Hardware	2018-01-02T00:00:00.000Z
2	All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi...	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	Electronics	2017-01-02T00:00:00.000Z
3	Brand New Amazon Kindle Fire 16gb 7" Ips Displ...	Amazon	Computers/Tablets & Networking,Tablets & eBook...	Electronics	2017-03-25T00:00:00.000Z
4	Amazon Echo Show Alexa-enabled Bluetooth Speake...	Amazon	Computers,Amazon Echo,Virtual Assistant Speake...	Electronics,Hardware	2017-11-15T00:00:00.000Z

In [6]:

```
train.count()
```

Out[6]:

```
name                4000
brand               4000
categories          4000
primaryCategories   4000
reviews.date        4000
reviews.text        4000
reviews.title       3990
sentiment           4000
dtype: int64
```

In [7]:

```
#train = train.append(test, ignore_index=True)
test.count()
```

Out[7]:

```
name                1000
brand               1000
categories          1000
primaryCategories   1000
reviews.date        1000
reviews.text        1000
reviews.title        997
dtype: int64
```

In [8]:

```
train.duplicated().sum()
```

Out[8]:

```
58
```

There are 58 duplicates, let's drop the duplicate values

In [9]:

```
train = train.drop_duplicates().reset_index(drop=True)
```

In [10]:

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3942 entries, 0 to 3941
Data columns (total 8 columns):
name                3942 non-null object
brand               3942 non-null object
categories          3942 non-null object
primaryCategories   3942 non-null object
reviews.date        3942 non-null object
reviews.text        3942 non-null object
reviews.title       3932 non-null object
sentiment           3942 non-null object
dtypes: object(8)
memory usage: 246.5+ KB
```

In [11]:

```
train.dtypes
```

Out[11]:

```
name                object
brand               object
categories          object
primaryCategories   object
reviews.date        object
reviews.text        object
reviews.title       object
sentiment           object
dtype: object
```

In [12]:

```
train.describe()
```

Out[12]:


	name	brand	categories	primaryCategories	reviews.date	reviews
count	3942	3942	3942	3942	3942	3942
unique	23	1	23	4	638	3598
top	Amazon Echo Show Alexa-enabled Bluetooth Speak...	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	Electronics	2017-01-23T00:00:00.000Z	I bough kindle f my 11y grandd
freq	676	3942	628	2562	98	4

Lets Visualize with the class imbalance thing !

Basic EDA of trainig data set


In [13]:

```
sns.countplot(y=train.sentiment);
```



In [14]:

```
sns.countplot( train['sentiment'] );
```



Class Imbalance Problem

In [15]:

```
train.sentiment.value_counts()
```

Out[15]:

Positive 3694

Neutral 158

Negative 90

Name: sentiment, dtype: int64

In [16]:

```
# NA data  
train.isnull().sum()
```

Out[16]:

```
name                0  
brand               0  
categories          0  
primaryCategories  0  
reviews.date       0  
reviews.text       0  
reviews.title      10  
sentiment          0  
dtype: int64
```

In [17]:

```
test.isnull().sum()
```

Out[17]:

```
name                0  
brand               0  
categories          0  
primaryCategories  0  
reviews.date       0  
reviews.text       0  
reviews.title       3  
dtype: int64
```

We should rename the column to avoid errors

In [18]:

```
train.columns
```

Out[18]:

```
Index(['name', 'brand', 'categories', 'primaryCategories', 'reviews.date',  
      'reviews.text', 'reviews.title', 'sentiment'],  
      dtype='object')
```

In [19]:

```
train.rename(columns = {'reviews.text':'reviews_text', 'reviews.title':'reviews_title',  
                        'reviews.date':'reviews_date'}, inplace = True)
```

In [20]:

```
train.columns
```

Out[20]:


```
Index(['name', 'brand', 'categories', 'primaryCategories', 'reviews_date',  
      'reviews_text', 'reviews_title', 'sentiment'],  
      dtype='object')
```

In [21]:

```
train['sentiment'].value_counts().plot(kind='pie', autopct= '%1.0f%%')
```

Out[21]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f720e0aa518>



Here I am dropping the Neutral Sentiment as My goal is to work on differentiate the positive and negative

In [22]:

```
train = train[train.sentiment != "Neutral"]
```

Now we are ready for a WordCloud visualization which shows only the most emphatic words of the Positive and Negative sentiment.

In [23]:

```
train_pos = train[ train['sentiment'] == 'Positive']
train_pos = train_pos['reviews_text']
train_neg = train[ train['sentiment'] == 'Negative']
train_neg = train_neg['reviews_text']

def wordcloud_draw(data, color = 'black'):
    words = ' '.join(data)
    cleaned_word = " ".join([word for word in words.split()
                              if 'http' not in word
                              and not word.startswith('@')
                              and not word.startswith('#')
                              and word != 'RT'
                              ])
    wordcloud = WordCloud(stopwords=STOPWORDS,
                          background_color=color,
                          width=2500,
                          height=2000
                          ).generate(cleaned_word)
    plt.figure(1,figsize=(13, 13))
    plt.imshow(wordcloud)
    plt.axis('off')
    plt.show()

print("Positive words")
wordcloud_draw(train_pos, 'white')
print("Negative words")
wordcloud_draw(train_neg)
```

Positive words



Negative words

In [24]:

```
def remove_non_ascii(words):
    """Remove non-ASCII characters from list of tokenized words"""
    new_words = []
    for word in words:
        new_word = unicodedata.normalize('NFKD', word).encode('ascii', 'ignore')
        .decode('utf-8', 'ignore')
        new_words.append(new_word)
    return new_words

def to_lowercase(words):
    """Convert all characters to lowercase from list of tokenized words"""
    new_words = []
    for word in words:
        new_word = word.lower()
        new_words.append(new_word)
    return new_words

def remove_punctuation(words):
    """Remove punctuation from list of tokenized words"""
    new_words = []
    for word in words:
        new_word = re.sub(r'[\W\s]', '', word)
        if new_word != '':
            new_words.append(new_word)
    return new_words

def remove_numbers(words):
    """Remove all interger occurrences in list of tokenized words with textual representation"""
    new_words = []
    for word in words:
        new_word = re.sub("\d+", "", word)
        if new_word != '':
            new_words.append(new_word)
    return new_words

def remove_stopwords(words):
    """Remove stop words from list of tokenized words"""
    new_words = []
    for word in words:
        if word not in stopwords.words('english'):
            new_words.append(word)
    return new_words
```

```

def stem_words(words):
    """Stem words in list of tokenized words"""
    stemmer = LancasterStemmer()
    stems = []
    for word in words:
        stem = stemmer.stem(word)
        stems.append(stem)
    return stems

def lemmatize_verbs(words):
    """Lemmatize verbs in list of tokenized words"""
    lemmatizer = WordNetLemmatizer()
    lemmas = []
    for word in words:
        lemma = lemmatizer.lemmatize(word, pos='v')
        lemmas.append(lemma)
    return lemmas

def normalize(words):
    words = remove_non_ascii(words)
    words = to_lowercase(words)
    words = remove_punctuation(words)
    words = remove_numbers(words)
    # words = remove_stopwords(words)
    return words

```

In [25]:

```

# First step - tokenizing phrases
train['reviews_text'] = train['reviews_text'].apply(nltk.word_tokenize)

# Second step - passing through prep functions
#train['reviews_text'] = train['reviews_text'].apply(normalize)
train['reviews_text'].head()

```

Out[25]:

```


0    [Purchased, on, Black, FridayPros, -, Great, P...
1    [I, purchased, two, Amazon, in, Echo, Plus, an...
3    [very, good, product, ., Exactly, what, I, wan...
4    [This, is, the, 3rd, one, I, 've, purchased, ....
5    [This, is, a, great, product, ., Light, weight...
Name: reviews_text, dtype: object

```

Let's visualize the most common keywords

In [26]:

```
# Most common keywords
plt.figure(figsize=(10,6))
sns.countplot(y=train.reviews_title, order = train.reviews_title.value_counts().
iloc[:20].index)
plt.title('Top 20 keywords')
plt.show()
# train.keyword.value_counts().head(10)
```



In []: